Measured compute, charged availability and the control path

Frozen CPU calibration is reused without retiming; actual wall times come from the completed confirmation request traces.

One request forecasts for all 100 robots. Startup is separate. Every family pays another 0.50 s FIFO delay.

A Warm requests: timing and charge

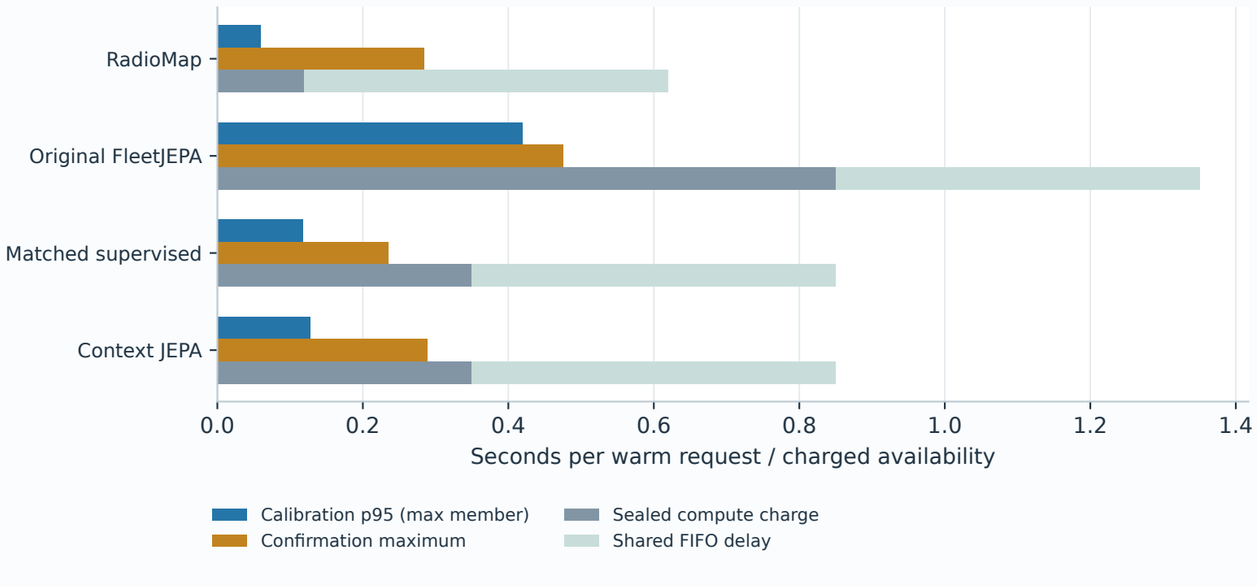

B Actual maxima and charges in every regime

Family	Cold	Refit	Warm
RadioMap	0.434 / 0.63	0.408 / 0.62	0.285 / 0.12
Original FleetJEPA	0.821 / 1.46	0.758 / 1.51	0.476 / 0.85
Matched supervised	0.816 / 1.63	0.494 / 0.93	0.235 / 0.35
Context JEPA	0.662 / 1.63	0.621 / 0.93	0.289 / 0.35

Cell = actual confirmation maximum / sealed compute charge (seconds).
Add 0.50 s FIFO delay to the charge for ready time; native slots round up.
New members share the original twelve-arm calibration envelope.
Global actual-overrun veto: TRIGGERED.

C Startup is separate from request compute

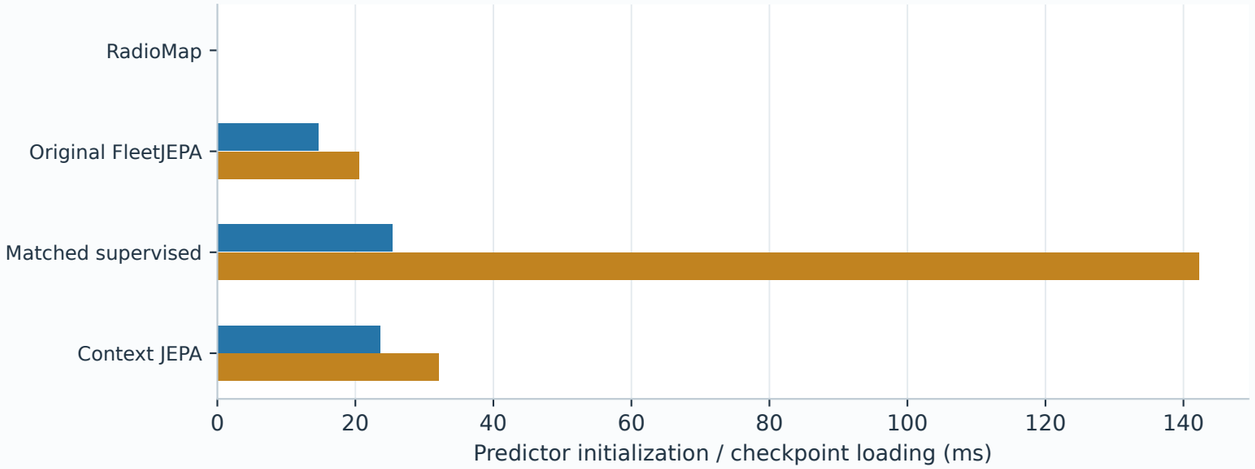

D Direct-horizon control path (schematic)

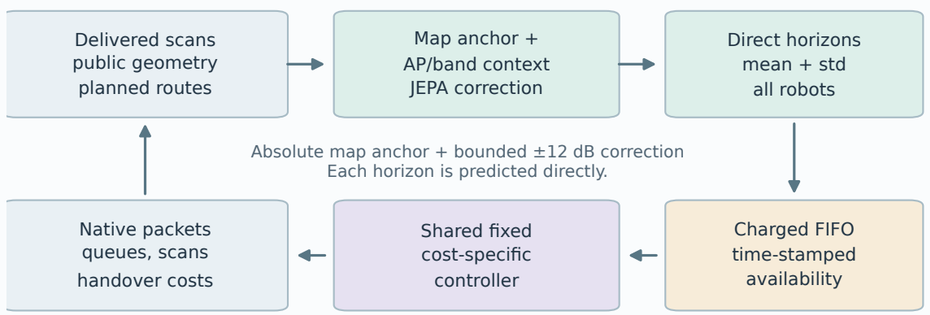

Timing summaries are descriptive maxima and calibration quantiles, not confidence intervals or device benchmarks. Calibration: one CPU thread, three repeated 24-step sequences per arm.
Confirmation maxima use 128 scientific rollouts; eight exact replays also enter the runtime veto. Predictor startup includes initialization/loading outside request charges.
Schematic only: the previously trained model uses simulator future-target supervision and an EMA context teacher, absent at inference. Original-ensemble comparisons cover the entire package.
Direct horizons avoid recursive summation of model corrections; they do not guarantee bounded cumulative packet loss. Simulation only; no physical-car or production qualification.